

Programming Assignment #3

Lottery Scheduler

2024. 11. 22

Overview

- Due date: 12/6
 - No late submission is allowed
- Objective
 - To gain further knowledge of a real kernel
 - To familiarize yourself with a scheduler
 - To change that scheduler to a new algorithm



Phase #1: Intro

- **Objective**
 - Your system call, `getreadcount()`, simply returns how many times that the `read()` system call has been called by user processes since the time that the kernel was booted.

Programming Environment

- **Install packages**

```
prompt> sudo apt update  
prompt> sudo apt install gcc make git qemu-system-x86 wget build-essential gdb
```

- **Download script for testing the code**

```
prompt> git clone https://github.com/Song-HyeonJin/HW3_xv6-test.git  
prompt> cd HW3_xv6-test/initial-xv6  
prompt> git clone https://github.com/Song-HyeonJin/xv6-public.git
```

- **You will do programming at “xv6-public” directory**



Programming Environment

- Test your code and submit the results

prompt> make qemu-nox

If all has worked well, you'll see this screen

```
SeaBIOS (version 1.15.0-1)

iPXE (https://ipxe.org) 00:03.0 CA00 PCI2.10 PnP PMM+1FF8B4A0+1FECB4A0 CA00

Booting from Hard Disk..xv6...
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
init: starting sh
$ |
```

Ctrl+a and press x to quit

```
prompt> cd ..../HW3_xv6-test/initial-xv6
prompt> ./test-getreadcount.sh
```



Phase #2: Scheduling-xv6-lottery

- **Objective**
 - Xv6 kernel does scheduling processes with round-robin
 - You should reimplement lottery scheduler in xv6 kernel
- **Implementations**
 - New system calls
 - Implementation int settickets(int number)
 - Implementation int getpinfo(struct pstat*)
 - Handle child process state
 - Modify fork() system call
 - Lottery scheduler
 - Reimplement scheduler() to a lottery scheduler algorithm
- **You will do programming at “xv6-public” directory**

Programming Environment

- Background
 - Ostep-projects
<https://github.com/remzi-arpacidusseau/ostep-projects/tree/master/scheduling-xv6-lottery>
 - xv6 book
<https://www.cs.virginia.edu/~cr4bd/4414/F2018/files/xv6book.pdf>
 - Lottery Scheduling
<http://pages.cs.wisc.edu/~remzi/OSTEP/cpu-sched-lottery.pdf>
- Follow the example of similar system calls (kill, fork, ⋯)
 - sysproc.c
 - syscall.c
 - syscall.h
 - usys.S
 - user.h
 - proc.c
 - defs.h
- You should pass arguments into kernel with **argptr()**, **argint()**
 - Good examples of how to pass arguments into the kernel are found in existing system calls.
Ex) sys_read() in sysfile.c



Program scheduler

- **int settickets(int number)**
 - Set the number of tickets of the calling process
 - This routine should return 0 if successful and -1 otherwise
Ex) if the caller passes in a number less than one, return -1.
- **int getpinfo(struct pstat *)**
 - Return some information about all running processes (pstat)
 - This routine should return 0 if successful and -1 otherwise
Ex) if a bad or NULL pointer is passed into the kernel, return -1

Program scheduler

- **Proc.h**

- Xv6 kernel stores “process state” in struct proc
- Each process has proc struct
- You’ll need to introduce some new variables in proc to implement lottery scheduler

```
39 struct proc {  
40     uint sz;                      // Size of process memory (bytes)  
41     pde_t* pgdir;                // Page table  
42     char *kstack;                // Bottom of kernel stack for this process  
43     enum procstate state;        // Process state  
44     int pid;                    // Process ID  
45     struct proc *parent;         // Parent process  
46     struct trapframe *tf;        // Trap frame for current syscall  
47     struct context *context;      // swtch() here to run process  
48     void *chan;                  // If non-zero, sleeping on chan  
49     int killed;                 // If non-zero, have been killed  
50     struct file *ofile[NOFILE];  // Open files  
51     struct inode *cwd;          // Current directory  
52     char name[16];              // Process name (debugging)  
53  
54     // For lottery scheduler  
55     int inuse;  
56     int ticks;  
57     int tickets;  
58 };  
59  
60 // Process memory is laid out contiguously, low addresses first:  
61 //   text  
62 //   original data and bss  
63 //   fixed-size stack  
64 //   expandable heap
```



Program scheduler

- **pstat.h**
 - You may need `struct pstat` for `int getpinfo(struct pstat*)`
 - If `getpinfo(struct pstat *p)` is invoked, you should store state of processes to `pstat` pointed by `p`

```
1 #ifndef _PSTAT_H_
2 #define _PSTAT_H_
3
4 #include "param.h"
5
6 struct pstat {
7     int inuse[NPROC];    // whether this slot of the process table is in use (1 or 0)
8     int tickets[NPROC]; // the number of tickets this process has
9     int pid[NPROC];     // the PID of each process
10    int ticks[NPROC];   // the number of ticks each process has accumulated
11 };
12
13#endif // _PSTAT_H_
```



Tickets of Child Process

- You'll need to make sure a child process inherits the same number of tickets as its parents
- Example:
 - If the parent has 10 tickets and calls fork() to create a child process
 - The child should also get 10 tickets
- You'll need to make sure a child tick is initialized
- Example
 - If the parent has 100 ticks, and calls fork() to create a child process
 - The child's ticks should be initialized to 0
- **Modify fork() in proc.c**

Lottery Scheduler

- Most of the code for the scheduler can be found in `proc.c`
- Modify the `scheduler()` in `proc.c`
 - Winner: a random number between [0, total # of tickets]
 - Choose a process to run based on tickets
- Use `rand()`, `srand()` in `proc.c`



Test Lottery Scheduler

- Use `xv6-public/test_lottery.c`
- `test_lottery.c`:
 - Create three processes (using `fork()`, `settickets()`)
 - Process A: 30 tickets
 - Process B: 20 tickets
 - Process C: 10 tickets
- You must implement “tickets of child process” for the test
 - `int setticket(int number)`
 - `int getpinfo`
- Show the number of ticks a set of three processes
 - Get `pstat.tick[NPROC]` using `getpinfo()`

Test Program

- You can implement your own test program
- Implement test program in xv6-public/ directory
 - ex) xv6-public/test.c
- Add make target for test program to Makefile
- Build and run xv6 kernel with ‘make qemu-nox’

```
SeaBIOS (version 1.15.0-1)

iPXE (https://ipxe.org) 00:03.0 CA00 PCI2.10 PnP PMM+1FF8B4A0+1FECB4A0 CA00

Booting from Hard Disk..xv6...
cpu0: starting 0
sb: size 1000 nblocks 941 minodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
init: starting sh
$ |
```

- ls : You can see your test program
- test: You can run your test program



Conclusion

- **Within Due: 12/6**
 - No late submission is allowed 
 - Take it step by step 